

Part1:

For this part, I just read the disparity map for each image and importing kitti_datahandler to read stereo calibration. Then I used the formula $z = f \cdot B / \text{disp}$ to get the depth map and assigned depth 0 to those with depth < 0.1 m and > 80 m. The result is in est_depth. Running command would be "python part1_estimate_depth.py test".

Part2:

After understanding and experimenting the parameter "confidence_th" and "threshold". I finally chose "confidence_th" = 0.1 and "threshold" = 0.6, which can recognize most objects accurately but don't have many duplicated bounding box for the same object identified. Having a lower "confidence_th" allows me to recognize more object at the beginning and then using proper threshold 0.6 to remove redundant bounding boxes for the same recognized object during non-maximum suppression. Running command is "python part2_yolo.py test".

Following are estimated bounding box in the test set

000011:



000012:



000013:



000014:



000015:



Part3:

After numerous experiments, I finally chose distance range as 10. Which means I will mask the pixels inside car bounding box with the depth in $(\text{average_depth}-10, \text{average_depth}+10)$. For the calculation of average_depth , I remove the 0 value's weights so it won't be affected by 0 depth a lot. The measurement of accuracy is precision and recall and also F score that combines those two together. The code of measurements is in `part3_segmentation.py`. In this way, it achieves overall above 80% precision and recall.

The result is in `est_seg`.

Running command is "python `part3_segmentation.py` test".